

DNN VS CNN

April 12, 2024

```
[1]: import zipfile
import os

# Path to the uploaded zip file
zip_path = 'archive (5).zip'
extract_folder = 'lfw_dataset'

# Extract the zip file
with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_folder)

# List the contents of the extracted folder
extracted_files = os.listdir(extract_folder)
extracted_files
```

```
[1]: ['lfw-deepfunneled',
      'lfw_allnames.csv',
      'lfw_readme.csv',
      'matchpairsDevTest.csv',
      'matchpairsDevTrain.csv',
      'mismatchpairsDevTest.csv',
      'mismatchpairsDevTrain.csv',
      'pairs.csv',
      'people.csv',
      'peopleDevTest.csv',
      'peopleDevTrain.csv']
```

```
[2]: from sklearn.datasets import fetch_lfw_people

# Load the Labeled Faces in the Wild dataset
lfw_people = fetch_lfw_people(min_faces_per_person=50, resize=0.5)

# Extract the features (images) and the target (labels)
X = lfw_people.data
y = lfw_people.target
lfw_people.images.shape, lfw_people.target_names
```

```
[2]: ((1560, 62, 47),
      array(['Ariel Sharon', 'Colin Powell', 'Donald Rumsfeld', 'George W Bush',
            'Gerhard Schroeder', 'Hugo Chavez', 'Jacques Chirac',
            'Jean Chretien', 'John Ashcroft', 'Junichiro Koizumi',
            'Serena Williams', 'Tony Blair'], dtype='<U17'))
```

```
[3]: from sklearn.datasets import fetch_lfw_people
      from sklearn.preprocessing import MinMaxScaler

      # Load the dataset with specified parameters
      lfw_people = fetch_lfw_people(min_faces_per_person=50, resize=0.5)
      X = lfw_people.data
      y = lfw_people.target

      # Apply Min-Max scaling
      scaler = MinMaxScaler()
      X_scaled = scaler.fit_transform(X)

      # Show the shape of the datasets
      X_scaled.shape, y.shape
```

```
[3]: ((1560, 2914), (1560,))
```

```
[4]: import tensorflow as tf
      from sklearn.model_selection import train_test_split
      from sklearn.datasets import fetch_lfw_people
      from sklearn.preprocessing import MinMaxScaler

      # Load LFW dataset
      lfw_dataset = fetch_lfw_people(min_faces_per_person=70, resize=0.4)
      X = lfw_dataset.data
      y = lfw_dataset.target

      # Apply min-max feature scaling
      scaler = MinMaxScaler()
      X_scaled = scaler.fit_transform(X)

      # Split dataset into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
      ↪random_state=42)

      # Define the model
      model = tf.keras.Sequential([
          tf.keras.layers.Dense(8, activation='relu', input_shape=(X_train.
      ↪shape[1],)),
          tf.keras.layers.Dense(16, activation='relu'),
          tf.keras.layers.Dense(32, activation='relu'),
```

```

tf.keras.layers.Dense(64, activation='relu'),
tf.keras.layers.Dense(128, activation='relu'),
tf.keras.layers.Dense(len(set(y)), activation='softmax')
])

# Compile the model
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(X_train, y_train, epochs=20, batch_size=64,
                    validation_split=0.2)

# Evaluate the model
test_loss, test_acc = model.evaluate(X_test, y_test)
print('Test accuracy:', test_acc)

```

C:\Users\kusum\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\dense.py:86: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```

Epoch 1/20
13/13          3s 29ms/step -
accuracy: 0.3026 - loss: 1.8545 - val_accuracy: 0.4029 - val_loss: 1.6756
Epoch 2/20
13/13          0s 7ms/step -
accuracy: 0.3943 - loss: 1.7111 - val_accuracy: 0.4029 - val_loss: 1.6541
Epoch 3/20
13/13          0s 8ms/step -
accuracy: 0.4066 - loss: 1.6909 - val_accuracy: 0.4223 - val_loss: 1.6906
Epoch 4/20
13/13          0s 7ms/step -
accuracy: 0.4172 - loss: 1.6961 - val_accuracy: 0.4175 - val_loss: 1.6379
Epoch 5/20
13/13          0s 7ms/step -
accuracy: 0.4409 - loss: 1.6206 - val_accuracy: 0.4320 - val_loss: 1.5616
Epoch 6/20
13/13          0s 7ms/step -
accuracy: 0.4686 - loss: 1.5373 - val_accuracy: 0.4757 - val_loss: 1.5683
Epoch 7/20
13/13          0s 7ms/step -
accuracy: 0.4724 - loss: 1.5074 - val_accuracy: 0.4806 - val_loss: 1.5320
Epoch 8/20
13/13          0s 6ms/step -
accuracy: 0.4756 - loss: 1.4329 - val_accuracy: 0.4951 - val_loss: 1.4681

```

```

Epoch 9/20
13/13          0s 7ms/step -
accuracy: 0.5082 - loss: 1.3516 - val_accuracy: 0.5097 - val_loss: 1.4584
Epoch 10/20
13/13          0s 7ms/step -
accuracy: 0.5277 - loss: 1.3274 - val_accuracy: 0.5146 - val_loss: 1.4526
Epoch 11/20
13/13          0s 7ms/step -
accuracy: 0.5036 - loss: 1.3675 - val_accuracy: 0.4757 - val_loss: 1.5819
Epoch 12/20
13/13          0s 6ms/step -
accuracy: 0.5049 - loss: 1.3343 - val_accuracy: 0.5049 - val_loss: 1.5086
Epoch 13/20
13/13          0s 6ms/step -
accuracy: 0.5192 - loss: 1.2711 - val_accuracy: 0.5388 - val_loss: 1.3994
Epoch 14/20
13/13          0s 6ms/step -
accuracy: 0.5624 - loss: 1.2309 - val_accuracy: 0.5194 - val_loss: 1.4175
Epoch 15/20
13/13          0s 7ms/step -
accuracy: 0.5566 - loss: 1.2173 - val_accuracy: 0.5243 - val_loss: 1.4054
Epoch 16/20
13/13          0s 8ms/step -
accuracy: 0.5326 - loss: 1.2955 - val_accuracy: 0.4903 - val_loss: 1.5770
Epoch 17/20
13/13          0s 7ms/step -
accuracy: 0.5026 - loss: 1.3340 - val_accuracy: 0.5146 - val_loss: 1.4017
Epoch 18/20
13/13          0s 7ms/step -
accuracy: 0.5417 - loss: 1.2733 - val_accuracy: 0.5146 - val_loss: 1.4918
Epoch 19/20
13/13          0s 7ms/step -
accuracy: 0.5649 - loss: 1.2117 - val_accuracy: 0.5194 - val_loss: 1.3823
Epoch 20/20
13/13          0s 7ms/step -
accuracy: 0.5587 - loss: 1.1989 - val_accuracy: 0.5388 - val_loss: 1.3446
9/9           0s 3ms/step -
accuracy: 0.5965 - loss: 1.1717
Test accuracy: 0.565891444683075

```

```

[33]: import tensorflow as tf
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
      from tensorflow.keras.optimizers import Adam

```

```

[44]: import numpy as np
      from sklearn.datasets import fetch_lfw_people

```

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping

model = Sequential([
    Input(shape=(62, 47, 1)),
    Conv2D(32, (3, 3), activation='relu', padding='same'),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu', padding='same'),
    MaxPooling2D((2, 2)),
    Conv2D(128, (3, 3), activation='relu', padding='same'),
    MaxPooling2D((2, 2)),
    Conv2D(128, (3, 3), activation='relu', padding='same'),
    Conv2D(256, (3, 3), activation='relu', padding='same'),
    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(y_categorical.shape[1], activation='softmax')
])

model.compile(optimizer=Adam(learning_rate=0.001),
↳loss='categorical_crossentropy', metrics=['accuracy'])

early_stopping = EarlyStopping(monitor='val_loss', patience=3,
↳restore_best_weights=True)

history = model.fit(
    X_train,
    y_train,
    epochs=30,
    batch_size=64,
    validation_data=(X_val, y_val),
    callbacks=[early_stopping]
)

test_loss, test_accuracy = model.evaluate(X_test, y_test)
print("Test Accuracy:", test_accuracy)

model.summary()

```

Epoch 1/30

20/20

5s 142ms/step -

accuracy: 0.1971 - loss: 2.5220 - val_accuracy: 0.3846 - val_loss: 2.0245
 Epoch 2/30
 20/20 5s 126ms/step -
 accuracy: 0.3444 - loss: 2.1579 - val_accuracy: 0.3846 - val_loss: 1.9650
 Epoch 3/30
 20/20 3s 153ms/step -
 accuracy: 0.3446 - loss: 2.1375 - val_accuracy: 0.3846 - val_loss: 2.0060
 Epoch 4/30
 20/20 3s 147ms/step -
 accuracy: 0.3470 - loss: 2.1566 - val_accuracy: 0.3846 - val_loss: 1.9843
 Epoch 5/30
 20/20 3s 147ms/step -
 accuracy: 0.3418 - loss: 2.1675 - val_accuracy: 0.3846 - val_loss: 1.9495
 Epoch 6/30
 20/20 3s 157ms/step -
 accuracy: 0.3346 - loss: 2.1480 - val_accuracy: 0.3846 - val_loss: 1.9380
 Epoch 7/30
 20/20 3s 154ms/step -
 accuracy: 0.3520 - loss: 2.0825 - val_accuracy: 0.3718 - val_loss: 1.9378
 Epoch 8/30
 20/20 5s 128ms/step -
 accuracy: 0.3565 - loss: 1.9992 - val_accuracy: 0.4167 - val_loss: 1.7605
 Epoch 9/30
 20/20 3s 150ms/step -
 accuracy: 0.4164 - loss: 1.8325 - val_accuracy: 0.5449 - val_loss: 1.5685
 Epoch 10/30
 20/20 3s 152ms/step -
 accuracy: 0.5015 - loss: 1.5140 - val_accuracy: 0.6282 - val_loss: 1.2682
 Epoch 11/30
 20/20 3s 143ms/step -
 accuracy: 0.6078 - loss: 1.2044 - val_accuracy: 0.5962 - val_loss: 1.2578
 Epoch 12/30
 20/20 3s 151ms/step -
 accuracy: 0.6373 - loss: 1.0935 - val_accuracy: 0.6795 - val_loss: 0.9722
 Epoch 13/30
 20/20 3s 143ms/step -
 accuracy: 0.7179 - loss: 0.8680 - val_accuracy: 0.7179 - val_loss: 0.8295
 Epoch 14/30
 20/20 3s 151ms/step -
 accuracy: 0.7568 - loss: 0.7545 - val_accuracy: 0.7308 - val_loss: 0.8416
 Epoch 15/30
 20/20 3s 152ms/step -
 accuracy: 0.7891 - loss: 0.6295 - val_accuracy: 0.7821 - val_loss: 0.7381
 Epoch 16/30
 20/20 3s 152ms/step -
 accuracy: 0.8212 - loss: 0.5427 - val_accuracy: 0.7885 - val_loss: 0.6402
 Epoch 17/30
 20/20 3s 148ms/step -

```

accuracy: 0.8495 - loss: 0.4380 - val_accuracy: 0.8141 - val_loss: 0.5804
Epoch 18/30
20/20          3s 152ms/step -
accuracy: 0.8861 - loss: 0.3297 - val_accuracy: 0.8397 - val_loss: 0.5960
Epoch 19/30
20/20          3s 152ms/step -
accuracy: 0.9021 - loss: 0.3135 - val_accuracy: 0.8462 - val_loss: 0.5143
Epoch 20/30
20/20          3s 154ms/step -
accuracy: 0.9366 - loss: 0.1877 - val_accuracy: 0.8397 - val_loss: 0.4603
Epoch 21/30
20/20          3s 150ms/step -
accuracy: 0.9530 - loss: 0.1503 - val_accuracy: 0.8718 - val_loss: 0.6142
Epoch 22/30
20/20          3s 149ms/step -
accuracy: 0.9512 - loss: 0.1690 - val_accuracy: 0.8590 - val_loss: 0.5277
Epoch 23/30
20/20          3s 150ms/step -
accuracy: 0.9614 - loss: 0.1308 - val_accuracy: 0.8782 - val_loss: 0.4879
5/5            0s 35ms/step -
accuracy: 0.8117 - loss: 0.6425
Test Accuracy: 0.8269230723381042

Model: "sequential_18"

```

Layer (type)	Output Shape	
Param #		
conv2d_85 (Conv2D)	(None, 62, 47, 32)	
↪ 320		
max_pooling2d_72 (MaxPooling2D)	(None, 31, 23, 32)	
↪ 0		
conv2d_86 (Conv2D)	(None, 31, 23, 64)	
↪ 18,496		
max_pooling2d_73 (MaxPooling2D)	(None, 15, 11, 64)	
↪ 0		
conv2d_87 (Conv2D)	(None, 15, 11, 128)	
↪ 73,856		
max_pooling2d_74 (MaxPooling2D)	(None, 7, 5, 128)	
↪ 0		

conv2d_88 (Conv2D)	(None, 7, 5, 128)	└
↳147,584		
conv2d_89 (Conv2D)	(None, 7, 5, 256)	└
↳295,168		
flatten_18 (Flatten)	(None, 8960)	└
↳ 0		
dense_67 (Dense)	(None, 512)	└
↳4,588,032		
dropout_2 (Dropout)	(None, 512)	└
↳ 0		
dense_68 (Dense)	(None, 12)	└
↳6,156		

Total params: 15,388,838 (58.70 MB)

Trainable params: 5,129,612 (19.57 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 10,259,226 (39.14 MB)

```
[45]: import numpy as np
from sklearn.datasets import fetch_lfw_people
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout, BatchNormalization, Activation
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.regularizers import l2

# Load the dataset
lfw_people = fetch_lfw_people(min_faces_per_person=50, resize=0.5)
X = lfw_people.data
y = lfw_people.target
```



```

# Apply Min-Max feature scaling
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

# Reshape X for CNN and convert y to categorical
X_resaped = X_scaled.reshape((-1, 62, 47, 1)) # Adjusting the dimensions for
↳CNN input
y_categorical = to_categorical(y)

# Split the data into training, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(X_resaped, y_categorical,
↳test_size=0.2, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5,
↳random_state=42)

# Data augmentation
datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)
datagen.fit(X_train)

# Build the CNN model with Batch Normalization and L2 Regularization
model = Sequential([
    Conv2D(32, (3, 3), padding='same', kernel_regularizer=l2(0.001),
↳input_shape=(62, 47, 1)),
    BatchNormalization(),
    Activation('relu'),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), padding='same', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    Activation('relu'),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), padding='same', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    Activation('relu'),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), padding='same', kernel_regularizer=l2(0.001)),

```

```

        BatchNormalization(),
        Activation('relu'),

        Conv2D(256, (3, 3), padding='same', kernel_regularizer=l2(0.001)),
        BatchNormalization(),
        Activation('relu'),

        Flatten(),
        Dense(512, activation='relu'),
        Dropout(0.5),
        Dense(y_categorical.shape[1], activation='softmax')
    ])

    # Compile the model
    model.compile(optimizer=Adam(learning_rate=0.0001),
                  loss='categorical_crossentropy', metrics=['accuracy'])

    # Set up early stopping
    early_stopping = EarlyStopping(monitor='val_loss', patience=5,
                                   restore_best_weights=True)

    # Train the model with data augmentation
    history = model.fit(
        datagen.flow(X_train, y_train, batch_size=64),
        epochs=50, # Increased epochs
        validation_data=(X_val, y_val),
        callbacks=[early_stopping]
    )

    # Evaluate the model on the test set
    test_loss, test_accuracy = model.evaluate(X_test, y_test)
    print("Test Accuracy:", test_accuracy)

    # Print the model architecture
    model.summary()

```

Epoch 1/50

C:\Users\kusum\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:99: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

    super().__init__(
C:\Users\kusum\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:120:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,

```

`max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

```
self._warn_if_super_not_called()
```

```
20/20          8s 243ms/step -  
accuracy: 0.2231 - loss: 3.2241 - val_accuracy: 0.1859 - val_loss: 2.8748  
Epoch 2/50  
20/20          5s 227ms/step -  
accuracy: 0.3219 - loss: 2.5316 - val_accuracy: 0.3846 - val_loss: 2.7810  
Epoch 3/50  
20/20          5s 228ms/step -  
accuracy: 0.3427 - loss: 2.4728 - val_accuracy: 0.3846 - val_loss: 2.7765  
Epoch 4/50  
20/20          5s 232ms/step -  
accuracy: 0.3781 - loss: 2.3632 - val_accuracy: 0.3846 - val_loss: 2.7484  
Epoch 5/50  
20/20          5s 231ms/step -  
accuracy: 0.3843 - loss: 2.3630 - val_accuracy: 0.3846 - val_loss: 2.7061  
Epoch 6/50  
20/20          5s 218ms/step -  
accuracy: 0.4214 - loss: 2.2317 - val_accuracy: 0.3846 - val_loss: 2.6199  
Epoch 7/50  
20/20          5s 227ms/step -  
accuracy: 0.4442 - loss: 2.2473 - val_accuracy: 0.4231 - val_loss: 2.6590  
Epoch 8/50  
20/20          5s 225ms/step -  
accuracy: 0.4464 - loss: 2.1765 - val_accuracy: 0.1410 - val_loss: 2.6388  
Epoch 9/50  
20/20          5s 232ms/step -  
accuracy: 0.4390 - loss: 2.1808 - val_accuracy: 0.3846 - val_loss: 2.5645  
Epoch 10/50  
20/20          5s 223ms/step -  
accuracy: 0.4740 - loss: 2.0563 - val_accuracy: 0.3846 - val_loss: 2.5460  
Epoch 11/50  
20/20          5s 238ms/step -  
accuracy: 0.4643 - loss: 2.0942 - val_accuracy: 0.3846 - val_loss: 2.4825  
Epoch 12/50  
20/20          5s 233ms/step -  
accuracy: 0.4968 - loss: 1.9854 - val_accuracy: 0.3846 - val_loss: 2.4511  
Epoch 13/50  
20/20          5s 226ms/step -  
accuracy: 0.4993 - loss: 1.9266 - val_accuracy: 0.3846 - val_loss: 2.5005  
Epoch 14/50  
20/20          5s 232ms/step -  
accuracy: 0.4965 - loss: 1.9621 - val_accuracy: 0.3846 - val_loss: 2.4461  
Epoch 15/50  
20/20          5s 240ms/step -  
accuracy: 0.5190 - loss: 1.8800 - val_accuracy: 0.3846 - val_loss: 2.4018
```

Epoch 16/50
20/20 5s 232ms/step -
accuracy: 0.5371 - loss: 1.8608 - val_accuracy: 0.3846 - val_loss: 2.3937

Epoch 17/50
20/20 5s 238ms/step -
accuracy: 0.5158 - loss: 1.8335 - val_accuracy: 0.3846 - val_loss: 2.2987

Epoch 18/50
20/20 5s 242ms/step -
accuracy: 0.5522 - loss: 1.7610 - val_accuracy: 0.3846 - val_loss: 2.3264

Epoch 19/50
20/20 5s 231ms/step -
accuracy: 0.5449 - loss: 1.8070 - val_accuracy: 0.3846 - val_loss: 2.3706

Epoch 20/50
20/20 5s 240ms/step -
accuracy: 0.5757 - loss: 1.7366 - val_accuracy: 0.3846 - val_loss: 2.3294

Epoch 21/50
20/20 5s 226ms/step -
accuracy: 0.5944 - loss: 1.6802 - val_accuracy: 0.3910 - val_loss: 2.1556

Epoch 22/50
20/20 5s 234ms/step -
accuracy: 0.5554 - loss: 1.7215 - val_accuracy: 0.3846 - val_loss: 2.2242

Epoch 23/50
20/20 5s 239ms/step -
accuracy: 0.5879 - loss: 1.6664 - val_accuracy: 0.4295 - val_loss: 2.0908

Epoch 24/50
20/20 5s 231ms/step -
accuracy: 0.5898 - loss: 1.6348 - val_accuracy: 0.5192 - val_loss: 1.8884

Epoch 25/50
20/20 5s 233ms/step -
accuracy: 0.6150 - loss: 1.5855 - val_accuracy: 0.4936 - val_loss: 1.8796

Epoch 26/50
20/20 5s 234ms/step -
accuracy: 0.6192 - loss: 1.5626 - val_accuracy: 0.6282 - val_loss: 1.5755

Epoch 27/50
20/20 5s 231ms/step -
accuracy: 0.6266 - loss: 1.5609 - val_accuracy: 0.6090 - val_loss: 1.5515

Epoch 28/50
20/20 5s 230ms/step -
accuracy: 0.6124 - loss: 1.5844 - val_accuracy: 0.6859 - val_loss: 1.3925

Epoch 29/50
20/20 5s 236ms/step -
accuracy: 0.6445 - loss: 1.5174 - val_accuracy: 0.7115 - val_loss: 1.3510

Epoch 30/50
20/20 5s 230ms/step -
accuracy: 0.6376 - loss: 1.5436 - val_accuracy: 0.6474 - val_loss: 1.3725

Epoch 31/50
20/20 5s 232ms/step -
accuracy: 0.6436 - loss: 1.5154 - val_accuracy: 0.7051 - val_loss: 1.3042

```

Epoch 32/50
20/20          5s 231ms/step -
accuracy: 0.6474 - loss: 1.4627 - val_accuracy: 0.7372 - val_loss: 1.1416
Epoch 33/50
20/20          5s 237ms/step -
accuracy: 0.6531 - loss: 1.4566 - val_accuracy: 0.7500 - val_loss: 1.1510
Epoch 34/50
20/20          5s 237ms/step -
accuracy: 0.6629 - loss: 1.4182 - val_accuracy: 0.7756 - val_loss: 1.0660
Epoch 35/50
20/20          5s 246ms/step -
accuracy: 0.6935 - loss: 1.3830 - val_accuracy: 0.7564 - val_loss: 1.1145
Epoch 36/50
20/20          5s 237ms/step -
accuracy: 0.6720 - loss: 1.3822 - val_accuracy: 0.7821 - val_loss: 1.0187
Epoch 37/50
20/20          5s 245ms/step -
accuracy: 0.6623 - loss: 1.3736 - val_accuracy: 0.7244 - val_loss: 1.1866
Epoch 38/50
20/20          5s 238ms/step -
accuracy: 0.6717 - loss: 1.3363 - val_accuracy: 0.7756 - val_loss: 1.0657
Epoch 39/50
20/20          5s 229ms/step -
accuracy: 0.7062 - loss: 1.3203 - val_accuracy: 0.7244 - val_loss: 1.1949
Epoch 40/50
20/20          5s 220ms/step -
accuracy: 0.6901 - loss: 1.3164 - val_accuracy: 0.7885 - val_loss: 1.1204
Epoch 41/50
20/20          5s 223ms/step -
accuracy: 0.6867 - loss: 1.3079 - val_accuracy: 0.7692 - val_loss: 1.1023
5/5           0s 25ms/step -
accuracy: 0.7059 - loss: 1.2397
Test Accuracy: 0.7243589758872986

Model: "sequential_19"

```

Layer (type)	Output Shape	
↳Param #		
conv2d_90 (Conv2D)	(None, 62, 47, 32)	↳
↳320		
batch_normalization	(None, 62, 47, 32)	↳
↳128		
(BatchNormalization)		↳
↳		

activation_5 (Activation)	(None, 62, 47, 32)	┐
↪ 0		
max_pooling2d_75 (MaxPooling2D)	(None, 31, 23, 32)	┐
↪ 0		
conv2d_91 (Conv2D)	(None, 31, 23, 64)	┐
↪18,496		
batch_normalization_1	(None, 31, 23, 64)	┐
↪256		
(BatchNormalization)		┐
↪		
activation_6 (Activation)	(None, 31, 23, 64)	┐
↪ 0		
max_pooling2d_76 (MaxPooling2D)	(None, 15, 11, 64)	┐
↪ 0		
conv2d_92 (Conv2D)	(None, 15, 11, 128)	┐
↪73,856		
batch_normalization_2	(None, 15, 11, 128)	┐
↪512		
(BatchNormalization)		┐
↪		
activation_7 (Activation)	(None, 15, 11, 128)	┐
↪ 0		
max_pooling2d_77 (MaxPooling2D)	(None, 7, 5, 128)	┐
↪ 0		
conv2d_93 (Conv2D)	(None, 7, 5, 128)	┐
↪147,584		
batch_normalization_3	(None, 7, 5, 128)	┐
↪512		
(BatchNormalization)		┐
↪		
activation_8 (Activation)	(None, 7, 5, 128)	┐
↪ 0		

```

conv2d_94 (Conv2D)                                (None, 7, 5, 256)
↳295,168
batch_normalization_4                            (None, 7, 5, 256)
↳1,024
(BatchNormalization)
↳
activation_9 (Activation)                        (None, 7, 5, 256)
↳ 0
flatten_19 (Flatten)                            (None, 8960)
↳ 0
dense_69 (Dense)                                (None, 512)
↳4,588,032
dropout_3 (Dropout)                             (None, 512)
↳ 0
dense_70 (Dense)                                (None, 12)
↳6,156

```

Total params: 15,393,702 (58.72 MB)

Trainable params: 5,130,828 (19.57 MB)

Non-trainable params: 1,216 (4.75 KB)

Optimizer params: 10,261,658 (39.15 MB)

```

[46]: # Define the accuracies obtained from the models
dnn_test_accuracy = 0.5659 # Update this with your DNN's test accuracy if it
↳changes
cnn_test_accuracy = 0.8269 # Update this with your CNN's test accuracy if it
↳changes

print("Classification Results:")
print("Deep Neural Network (DNN) Test Accuracy: {:.2f}%".
↳format(dnn_test_accuracy * 100))
print("Convolutional Neural Network (CNN) Test Accuracy: {:.2f}%".
↳format(cnn_test_accuracy * 100))

```

Classification Results:

Deep Neural Network (DNN) Test Accuracy: 56.59%

Convolutional Neural Network (CNN) Test Accuracy: 82.69%

Report: Classification Results and Analysis
Classification Results: Deep Neural Network (DNN) Test Accuracy: 56.59% Convolutional Neural Network (CNN) Test Accuracy: 82.69% These results indicate that the CNN significantly outperformed the DNN in this classification task.

Analysis Based on Classification Results a. How do the learning rate and filter size value affect the CNN classification accuracy?

The learning rate and filter size are fundamental hyperparameters in training convolutional neural networks:

Learning Rate: This controls the step size during the gradient descent optimization process. An optimally set learning rate helps the model converge efficiently to a good solution. If the learning rate is too high, the model may converge too quickly to a suboptimal solution or overshoot the minima. If it's too low, the model may take too long to converge, leading to lengthy training times or getting stuck in local minima. The CNN's higher accuracy suggests that the learning rate was well-tuned for balancing the speed of convergence with the stability of the training process.

Filter Size: This determines the size of the kernel that moves over the input image to extract features. Smaller filters might capture finer details, while larger filters tend to capture broader features of the image. The choice of filter size affects the model's ability to generalize from the training data to unseen data. In this case, the effective use of filter sizes likely contributed to the CNN's ability to recognize and generalize critical features from the input data, resulting in higher accuracy.

b. Does the increased hidden layer in DNN size improve the classification accuracy?

Increasing the size or number of hidden layers in a neural network can improve its capacity to learn more complex features from the data, which is often beneficial for tasks involving high-dimensional input such as images. However, there are several caveats:

Capacity vs. Overfitting: While more layers can increase the model's capacity to learn, they can also make the model prone to overfitting, especially if there isn't enough data to support the complexity of the model. Overfitting results in high performance on the training data but poor generalization to new, unseen data.

Training Challenges: Larger networks are more challenging to train. They are more susceptible to issues like vanishing gradients, where gradients become too small for effective training as they propagate back through the layers during training. This can be mitigated with techniques like proper initialization, using ReLU activations, or employing batch normalization.

In this scenario, despite increasing the DNN's hidden layer size, the CNN still outperformed it. This outcome suggests that merely increasing the size of the network was not sufficient to achieve better results compared to the CNN, which efficiently leverages spatial hierarchies and parameter sharing through convolutional layers. It's possible that the DNN could improve with more data, better regularization strategies, or an optimized architecture specifically suited for the dataset in use.

Conclusion: The comparison clearly demonstrates the advantages of using CNNs for image-based tasks over DNNs due to their architectural benefits in handling spatial data. Adjusting learning

rates and filter sizes critically influences CNN performance, while simply increasing the DNN size without strategic considerations can lead to inefficiencies and suboptimal learning.